

complex a DIS final draft V1

August 28, 2025

1 1. Imports & Setup

```
[5]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix
import torch
import torch.nn as torch.nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
import shap
import lime
import lime.lime_tabular
from lime.lime_tabular import LimeTabularExplainer
```

```
[6]: np.random.seed(42)
torch.manual_seed(42)
```

```
[6]: <torch._C.Generator at 0x32014a6b0>
```

2 2. Data Exploartion & Preprocessing

```
[7]: df = pd.read_csv('A_50000.csv')
print(f"Dataset shape: {df.shape}")
print(f"Target distribution: {df['eligible'].value_counts().to_dict()}")

# Identify feature types
feature_cols = [col for col in df.columns if col != 'eligible']
noise_features = [col for col in feature_cols if col.startswith('noise_')]
real_features = [col for col in feature_cols if not col.startswith('noise_')]
```

```

print(f"Total features: {len(feature_cols)}")
print(f"Noise features: {len(noise_features)}")
print(f"Real features: {len(real_features)}")
print(f"Real features: {real_features}")

# Encode categorical variables
df_processed = df.copy()

categorical_cols = ['gender', 'patient_type']
for col in df.columns:
    if col.startswith('paid_contribution_'):
        categorical_cols.append(col)

for col in categorical_cols:
    if col in df_processed.columns:
        le = LabelEncoder()
        df_processed[col] = le.fit_transform(df_processed[col])

# Convert boolean columns
boolean_cols = ['is_spouse', 'is_absent', 'eligible']
for col in boolean_cols:
    if col in df_processed.columns:
        df_processed[col] = df_processed[col].astype(int)

# Prepare final dataset
X = df_processed[feature_cols]
y = df_processed['eligible']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print(f"Training samples: {len(X_train)}")
print(f"Test samples: {len(X_test)}")

```

Dataset shape: (50000, 65)

Target distribution: {True: 25000, False: 25000}

Total features: 64

Noise features: 52

Real features: 12

Real features: ['age', 'gender', 'paid_contribution_1', 'paid_contribution_2', 'paid_contribution_3', 'paid_contribution_4', 'paid_contribution_5',

```
'is_spouse', 'is_absent', 'capital_resources', 'patient_type',  
'distance_to_hospital']  
Training samples: 40000  
Test samples: 10000
```

3 3. Training Traditional Models

```
[8]: # Logistic Regression  
print("Training Logistic Regression...")  
lr = LogisticRegression(random_state=42, max_iter=1000, C=0.1)  
lr.fit(X_train_scaled, y_train)  
lr_pred = lr.predict(X_test_scaled)  
lr_acc = accuracy_score(y_test, lr_pred)  
print(f"Logistic Regression accuracy: {lr_acc:.4f}")  
  
# Random Forest  
print("Training Random Forest...")  
rf = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=42)  
rf.fit(X_train, y_train)  
rf_pred = rf.predict(X_test)  
rf_acc = accuracy_score(y_test, rf_pred)  
print(f"Random Forest accuracy: {rf_acc:.4f}")  
  
# Neural Network  
print("Training Neural Network...")  
nn = MLPClassifier(hidden_layer_sizes=(100, 50), random_state=42, max_iter=500,  
    ↪alpha=0.01)  
nn.fit(X_train_scaled, y_train)  
nn_pred = nn.predict(X_test_scaled)  
nn_acc = accuracy_score(y_test, nn_pred)  
print(f"Neural Network accuracy: {nn_acc:.4f}")
```

```
Training Logistic Regression...  
Logistic Regression accuracy: 0.9939  
Training Random Forest...  
Random Forest accuracy: 0.9981  
Training Neural Network...  
Neural Network accuracy: 0.9973
```

4 4. KAN Implementation

```
[9]: class KAN(torch.nn.Module):  
    def __init__(self, input_dim, hidden_dim, output_dim, grid_size=5):  
        super(KAN, self).__init__()  
        self.input_dim = input_dim  
        self.hidden_dim = hidden_dim
```

```

self.grid_size = grid_size

# Learnable spline coefficients
self.spline_coeffs = torch.nn.Parameter(
    torch.randn(input_dim, hidden_dim, grid_size) * 0.1
)

# Connection scales
self.scales = torch.nn.Parameter(torch.ones(input_dim, hidden_dim) * 0.1)
↪1)

# Output layer
self.output_layer = torch.nn.Linear(hidden_dim, output_dim)
self.dropout = torch.nn.Dropout(0.1)

def spline_activation(self, x, coeffs):
    x_norm = torch.sigmoid(x)
    batch_size = x_norm.shape[0]
    grid_points = torch.linspace(0, 1, self.grid_size, device=x.device)

    basis_values = []
    for i in range(self.grid_size):
        distances = torch.abs(x_norm - grid_points[i])
        basis = torch.clamp(1.0 - distances * self.grid_size, 0.0, 1.0)
        basis_values.append(basis)

    basis_tensor = torch.stack(basis_values, dim=-1)
    weighted_basis = basis_tensor * coeffs.unsqueeze(0)
    return torch.sum(weighted_basis, dim=-1)

def forward(self, x):
    batch_size = x.shape[0]
    hidden_outputs = []

    for h in range(self.hidden_dim):
        hidden_val = 0
        for i in range(self.input_dim):
            x_input = x[:, i:i+1]
            spline_out = self.spline_activation(x_input, self.
↪spline_coeffs[i, h, :])
            scaled_out = spline_out * self.scales[i, h]
            hidden_val += scaled_out
        hidden_outputs.append(hidden_val)

    hidden = torch.cat(hidden_outputs, dim=1)
    hidden = torch.tanh(hidden)
    hidden = self.dropout(hidden)

```

```

        output = self.output_layer(hidden)
        return output

```

5 5. KAN Training

```

[10]: # Data preparation
X_train_tensor = torch.FloatTensor(X_train_scaled)
X_test_tensor = torch.FloatTensor(X_test_scaled)
y_train_tensor = torch.FloatTensor(y_train.values)
y_test_tensor = torch.FloatTensor(y_test.values)

train_dataset = TensorDataset(X_train_tensor, y_train_tensor)
train_loader = DataLoader(train_dataset, batch_size=512, shuffle=True)

# Model initialization
input_size = X_train_scaled.shape[1]
kan_model = KAN(input_dim=input_size, hidden_dim=8, output_dim=1, grid_size=7)

# Training setup
optimizer = optim.Adam(kan_model.parameters(), lr=0.001, weight_decay=1e-4)
criterion = torch.nn.BCEWithLogitsLoss()

print(f"Model parameters: {sum(p.numel() for p in kan_model.parameters())}")

# Training loop
kan_model.train()
for epoch in range(100):
    epoch_loss = 0
    num_batches = 0

    for batch_x, batch_y in train_loader:
        optimizer.zero_grad()
        outputs = kan_model(batch_x).squeeze()
        loss = criterion(outputs, batch_y)

        # L1 regularization
        l1_reg = torch.sum(torch.abs(kan_model.spline_coeffs))
        total_loss = loss + 0.001 * l1_reg

        total_loss.backward()
        optimizer.step()

        epoch_loss += total_loss.item()
        num_batches += 1

    if epoch % 25 == 0:

```

```

        avg_loss = epoch_loss / num_batches
        print(f"Epoch {epoch}: Loss = {avg_loss:.4f}")

# KAN evaluation
kan_model.eval()
with torch.no_grad():
    kan_outputs = kan_model(X_test_tensor).squeeze()
    kan_probs = torch.sigmoid(kan_outputs)
    kan_pred = (kan_probs > 0.5).float().numpy()
    kan_acc = accuracy_score(y_test, kan_pred)

```

Model parameters: 4105
 Epoch 0: Loss = 0.8525
 Epoch 25: Loss = 0.0623
 Epoch 50: Loss = 0.0403
 Epoch 75: Loss = 0.0299

6 6. Results Comparison

```

[11]: results = {
        'Logistic Regression': lr_acc,
        'Random Forest': rf_acc,
        'Neural Network': nn_acc,
        'KAN': kan_acc
    }

    for model, accuracy in results.items():
        print(f"{model}: {accuracy:.4f}")

```

Logistic Regression: 0.9939
 Random Forest: 0.9981
 Neural Network: 0.9973
 KAN: 0.9970

7 7. Feature Importance

```

[16]: # RF Features importance
rf_importance = rf.feature_importances_
rf_feature_ranking = list(zip(feature_cols, rf_importance))

# Logistic Regression importance
lr_importance = np.abs(lr.coef_[0])
lr_feature_ranking = list(zip(feature_cols, lr_importance))

# Your KAN importance
def get_kan_importance():

```

```

kan_model.eval()
baseline_input = torch.FloatTensor(X_test_scaled.mean(axis=0)).unsqueeze(0)

with torch.no_grad():
    baseline_output = kan_model(baseline_input).item()

sensitivities = []
for i in range(len(feature_cols)):
    feature_values = np.percentile(X_test_scaled[:, i], [10, 50, 90])
    outputs = []

    for val in feature_values:
        test_input = baseline_input.clone()
        test_input[0, i] = val
        with torch.no_grad():
            output = kan_model(test_input).item()
        outputs.append(output)

    sensitivity = np.max(outputs) - np.min(outputs)
    sensitivities.append(sensitivity)

return np.array(sensitivities)

# KAN importance calculation
kan_importance = get_kan_importance()
kan_feature_ranking = list(zip(feature_cols, kan_importance))

# Display results
print("Enhanced Feature Importance Analysis:")

print("\nLogistic Regression - Top 10:")
lr_sorted = sorted(lr_feature_ranking, key=lambda x: x[1], reverse=True)
for i, (feature, importance) in enumerate(lr_sorted[:10], 1):
    feature_type = "REAL" if feature in real_features else "NOISE"
    print(f"{i:2d}. {feature:25s}: {importance:.4f} ({feature_type})")

print(f"\nRandom Forest - Top 10:")
rf_sorted = sorted(rf_feature_ranking, key=lambda x: x[1], reverse=True)
for i, (feature, importance) in enumerate(rf_sorted[:10], 1):
    feature_type = "REAL" if feature in real_features else "NOISE"
    print(f"{i:2d}. {feature:25s}: {importance:.4f} ({feature_type})")

print(f"\nKAN - Top 10:")
kan_sorted = sorted(kan_feature_ranking, key=lambda x: x[1], reverse=True)
for i, (feature, importance) in enumerate(kan_sorted[:10], 1):
    feature_type = "REAL" if feature in real_features else "NOISE"
    print(f"{i:2d}. {feature:25s}: {importance:.4f} ({feature_type})")

```

Enhanced Feature Importance Analysis:

Logistic Regression - Top 10:

1. is_absent	: 3.3369 (REAL)
2. is_spouse	: 3.3199 (REAL)
3. capital_resources	: 3.1139 (REAL)
4. age	: 2.4009 (REAL)
5. paid_contribution_2	: 1.6521 (REAL)
6. paid_contribution_4	: 1.6455 (REAL)
7. paid_contribution_1	: 1.6233 (REAL)
8. paid_contribution_5	: 1.5841 (REAL)
9. paid_contribution_3	: 1.5672 (REAL)
10. gender	: 0.1622 (REAL)

Random Forest - Top 10:

1. capital_resources	: 0.2503 (REAL)
2. age	: 0.2146 (REAL)
3. is_absent	: 0.1413 (REAL)
4. is_spouse	: 0.1399 (REAL)
5. paid_contribution_2	: 0.0542 (REAL)
6. paid_contribution_1	: 0.0541 (REAL)
7. paid_contribution_5	: 0.0534 (REAL)
8. paid_contribution_3	: 0.0402 (REAL)
9. paid_contribution_4	: 0.0390 (REAL)
10. patient_type	: 0.0006 (REAL)

KAN - Top 10:

1. capital_resources	: 8.2440 (REAL)
2. age	: 8.2002 (REAL)
3. is_absent	: 8.1590 (REAL)
4. is_spouse	: 8.1133 (REAL)
5. paid_contribution_2	: 3.8879 (REAL)
6. paid_contribution_3	: 3.7750 (REAL)
7. paid_contribution_1	: 3.7309 (REAL)
8. paid_contribution_5	: 3.5878 (REAL)
9. paid_contribution_4	: 3.5196 (REAL)
10. noise_40	: 0.0000 (NOISE)

8 8. Classification Reports

```
[17]: print("Logistic Regression:")
print(classification_report(y_test, lr_pred, target_names=['Not Eligible',
↳ 'Eligible']))

print("Random Forest:")
print(classification_report(y_test, rf_pred, target_names=['Not Eligible',
↳ 'Eligible']))
```



```

print("Neural Network:")
print(classification_report(y_test, nn_pred, target_names=['Not Eligible',
↳ 'Eligible']))

print("KAN:")
print(classification_report(y_test, kan_pred, target_names=['Not Eligible',
↳ 'Eligible']))

```

Logistic Regression:

	precision	recall	f1-score	support
Not Eligible	1.00	0.99	0.99	5000
Eligible	0.99	1.00	0.99	5000
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

Random Forest:

	precision	recall	f1-score	support
Not Eligible	1.00	1.00	1.00	5000
Eligible	1.00	1.00	1.00	5000
accuracy			1.00	10000
macro avg	1.00	1.00	1.00	10000
weighted avg	1.00	1.00	1.00	10000

Neural Network:

	precision	recall	f1-score	support
Not Eligible	1.00	0.99	1.00	5000
Eligible	1.00	1.00	1.00	5000
accuracy			1.00	10000
macro avg	1.00	1.00	1.00	10000
weighted avg	1.00	1.00	1.00	10000

KAN:

	precision	recall	f1-score	support
Not Eligible	1.00	0.99	1.00	5000
Eligible	0.99	1.00	1.00	5000
accuracy			1.00	10000
macro avg	1.00	1.00	1.00	10000
weighted avg	1.00	1.00	1.00	10000

9. Rule Based Testing

```
[20]: def create_comprehensive_rule_tests():  
    """Create comprehensive test cases for all rules without manipulation"""  
    test_cases = []  
  
    # Base case that satisfies ALL rules  
    base_case = {  
        'age': 67,  
        'gender': 1, # Male  
        'is_spouse': 1,  
        'is_absent': 0,  
        'capital_resources': 2000,  
        'patient_type': 1, # In-patient  
        'distance_to_hospital': 30,  
        'paid_contribution_1': 1,  
        'paid_contribution_2': 1,  
        'paid_contribution_3': 1,  
        'paid_contribution_4': 1,  
        'paid_contribution_5': 1  
    }  
  
    # Add noise features to base case  
    for i in range(1, 53):  
        base_case[f'noise_{i}'] = X_train[f'noise_{i}'].mean()  
  
    # Test Rule 1: Age-Gender thresholds  
    test_cases.extend([  
        ('male_64_should_fail', {**base_case, 'age': 64, 'gender': 1}, False),  
        ('male_65_should_pass', {**base_case, 'age': 65, 'gender': 1}, True),  
        ('male_66_should_pass', {**base_case, 'age': 66, 'gender': 1}, True),  
        ('female_59_should_fail', {**base_case, 'age': 59, 'gender': 0}, False),  
        ('female_60_should_pass', {**base_case, 'age': 60, 'gender': 0}, True),  
        ('female_61_should_pass', {**base_case, 'age': 61, 'gender': 0}, True),  
    ])  
  
    # Test Rule 2: Contribution requirements (at least 4 of 5)  
    test_cases.extend([  
        ('contrib_5of5_pass', {**base_case}, True),  
        ('contrib_4of5_pass', {**base_case, 'paid_contribution_5': 0}, True),  
        ('contrib_3of5_fail', {**base_case, 'paid_contribution_4': 0,   
↪ 'paid_contribution_5': 0}, False),  
        ('contrib_2of5_fail', {**base_case, 'paid_contribution_3': 0,   
↪ 'paid_contribution_4': 0, 'paid_contribution_5': 0}, False),
```

```

])

# Test Rule 3: Spouse requirement
test_cases.extend([
    ('spouse_yes_pass', {**base_case}, True),
    ('spouse_no_fail', {**base_case, 'is_spouse': 0}, False),
])

# Test Rule 4: Absence requirement
test_cases.extend([
    ('absent_no_pass', {**base_case}, True),
    ('absent_yes_fail', {**base_case, 'is_absent': 1}, False),
])

# Test Rule 5: Capital resources limit
test_cases.extend([
    ('capital_2000_pass', {**base_case}, True),
    ('capital_3000_pass', {**base_case, 'capital_resources': 3000}, True),
    ('capital_3001_fail', {**base_case, 'capital_resources': 3001}, False),
    ('capital_5000_fail', {**base_case, 'capital_resources': 5000}, False),
])

# Test Rule 6: Patient type and distance
test_cases.extend([
    ('inpatient_30km_pass', {**base_case, 'patient_type': 1,
↪ 'distance_to_hospital': 30}, True),
    ('inpatient_50km_pass', {**base_case, 'patient_type': 1,
↪ 'distance_to_hospital': 50}, True),
    ('inpatient_51km_fail', {**base_case, 'patient_type': 1,
↪ 'distance_to_hospital': 51}, False),
    ('outpatient_51km_pass', {**base_case, 'patient_type': 0,
↪ 'distance_to_hospital': 51}, True),
    ('outpatient_50km_fail', {**base_case, 'patient_type': 0,
↪ 'distance_to_hospital': 50}, False),
    ('outpatient_30km_fail', {**base_case, 'patient_type': 0,
↪ 'distance_to_hospital': 30}, False),
])

# Complex combinations (multiple rule violations)
test_cases.extend([
    ('young_male_few_contrib', {**base_case, 'age': 60,
↪ 'paid_contribution_5': 0}, False), # Fails age AND contribution
    ('old_female_not_spouse', {**base_case, 'age': 65, 'gender': 0,
↪ 'is_spouse': 0}, False), # Passes age but fails spouse

```

```

        ('edge_case_all_boundaries', {**base_case, 'age': 65, 'gender': 1,
↪ 'capital_resources': 3000, 'paid_contribution_5': 0}, True), # All at
↪ boundaries
    ])

    return test_cases

def evaluate_models_on_rules(test_cases):
    """Evaluate all models on rule-based test cases"""
    results = {
        'LR': {'correct': 0, 'total': 0, 'by_rule': {}},
        'RF': {'correct': 0, 'total': 0, 'by_rule': {}},
        'NN': {'correct': 0, 'total': 0, 'by_rule': {}},
        'KAN': {'correct': 0, 'total': 0, 'by_rule': {}}
    }

    detailed_results = []

    for test_name, test_data, expected in test_cases:
        # Prepare test data
        test_df = pd.DataFrame([test_data])
        test_X = test_df[feature_cols]
        test_X_scaled = scaler.transform(test_X)

        # Get predictions
        lr_prob = lr.predict_proba(test_X_scaled)[0, 1]
        rf_prob = rf.predict_proba(test_X)[0, 1]
        nn_prob = nn.predict_proba(test_X_scaled)[0, 1]

        kan_model.eval()
        with torch.no_grad():
            kan_logit = kan_model(torch.FloatTensor(test_X_scaled)).item()
            kan_prob = torch.sigmoid(torch.tensor(kan_logit)).item()

        # Check predictions
        lr_pred = lr_prob > 0.5
        rf_pred = rf_prob > 0.5
        nn_pred = nn_prob > 0.5
        kan_pred = kan_prob > 0.5

        # Store detailed results
        detailed_results.append({
            'test': test_name,
            'expected': expected,
            'LR_prob': lr_prob,
            'RF_prob': rf_prob,

```

```

        'NN_prob': nn_prob,
        'KAN_prob': kan_prob,
        'LR_correct': lr_pred == expected,
        'RF_correct': rf_pred == expected,
        'NN_correct': nn_pred == expected,
        'KAN_correct': kan_pred == expected
    })

    # Update summary stats
    for model in ['LR', 'RF', 'NN', 'KAN']:
        model_correct = detailed_results[-1][f'{model}_correct']
        results[model]['correct'] += model_correct
        results[model]['total'] += 1

        # Track by rule category
        rule_category = test_name.split('_')[0]
        if rule_category not in results[model]['by_rule']:
            results[model]['by_rule'][rule_category] = {'correct': 0,
↳ 'total': 0}
            results[model]['by_rule'][rule_category]['correct'] += model_correct
            results[model]['by_rule'][rule_category]['total'] += 1

    return results, detailed_results

def plot_rule_test_results(results, detailed_results):
    """Visualize model performance on rule tests"""
    fig, axes = plt.subplots(2, 3, figsize=(18, 10))

    # 1. Overall accuracy
    models = ['LR', 'RF', 'NN', 'KAN']
    accuracies = [results[m]['correct']/results[m]['total'] for m in models]
    colors = ['blue', 'green', 'orange', 'red']

    axes[0, 0].bar(models, accuracies, color=colors, alpha=0.7)
    axes[0, 0].set_title('Overall Rule Test Accuracy')
    axes[0, 0].set_ylabel('Accuracy')
    axes[0, 0].set_ylim(0, 1)

    for i, acc in enumerate(accuracies):
        axes[0, 0].text(i, acc + 0.02, f'{acc:.2%}', ha='center')

    # 2. Performance by rule category
    rule_categories = ['male', 'female', 'contrib', 'spouse', 'absent',
↳ 'capital', 'inpatient', 'outpatient']
    category_performance = {model: [] for model in models}

```

```

for category in rule_categories:
    for model in models:
        if category in results[model]['by_rule']:
            acc = results[model]['by_rule'][category]['correct'] /
↳ results[model]['by_rule'][category]['total']
            category_performance[model].append(acc)
        else:
            category_performance[model].append(0)

x = np.arange(len(rule_categories))
width = 0.2

for i, model in enumerate(models):
    axes[0, 1].bar(x + i*width, category_performance[model], width,
↳ label=model, color=colors[i], alpha=0.7)

axes[0, 1].set_xlabel('Rule Category')
axes[0, 1].set_ylabel('Accuracy')
axes[0, 1].set_title('Performance by Rule Category')
axes[0, 1].set_xticks(x + width * 1.5)
axes[0, 1].set_xticklabels(rule_categories, rotation=45)
axes[0, 1].legend()
axes[0, 1].grid(True, alpha=0.3)

# 3. Confusion matrix for each model
for idx, model in enumerate(models):
    row = 0 if idx < 2 else 1
    col = 2 if idx % 2 == 0 else 0
    if row == 1:
        col = idx - 2

    # Calculate confusion matrix
    tp = sum(1 for r in detailed_results if r['expected'] and
↳ r[f'{model}_correct'])
    tn = sum(1 for r in detailed_results if not r['expected'] and
↳ r[f'{model}_correct'])
    fp = sum(1 for r in detailed_results if not r['expected'] and not
↳ r[f'{model}_correct'])
    fn = sum(1 for r in detailed_results if r['expected'] and not
↳ r[f'{model}_correct'])

    cm = np.array([[tn, fp], [fn, tp]])
    im = axes[1, col].imshow(cm, cmap='Blues')
    axes[1, col].set_title(f'{model} Confusion Matrix')
    axes[1, col].set_xlabel('Predicted')
    axes[1, col].set_ylabel('Actual')

```

```

axes[1, col].set_xticks([0, 1])
axes[1, col].set_yticks([0, 1])
axes[1, col].set_xticklabels(['Ineligible', 'Eligible'])
axes[1, col].set_yticklabels(['Ineligible', 'Eligible'])

# Add text annotations
for i in range(2):
    for j in range(2):
        axes[1, col].text(j, i, str(cm[i, j]), ha='center',
↪va='center', color='black', fontsize=12)

# Hide the last subplot
axes[1, 2].axis('off')

plt.tight_layout()
plt.savefig('rule_test_performance.png', dpi=300, bbox_inches='tight')
plt.show()

def print_detailed_failure_analysis(detailed_results):
    print("DETAILED FAILURE ANALYSIS")

    models = ['LR', 'RF', 'NN', 'KAN']

    for model in models:
        failures = [r for r in detailed_results if not r[f'{model}_correct']]
        print(f"\n{n{model} Failures ({len(failures)} total):")
        print("-" * 40)

        for failure in failures[:10]: # Show first 10 failures
            exp_str = "ELIGIBLE" if failure['expected'] else "INELIGIBLE"
            pred_prob = failure[f'{model}_prob']
            pred_str = "ELIGIBLE" if pred_prob > 0.5 else "INELIGIBLE"
            print(f"  {failure['test']:30s}: Expected {exp_str:10s}, Got
↪{pred_str:10s} (prob={pred_prob:.3f})")

        if len(failures) > 10:
            print(f"    ... and {len(failures)-10} more failures")

# Run comprehensive rule testing
print("COMPREHENSIVE RULE TESTING")

# Create test cases
test_cases = create_comprehensive_rule_tests()
print(f"\nCreated {len(test_cases)} test cases covering all rules")

```

```

# Evaluate models
results, detailed_results = evaluate_models_on_rules(test_cases)

# Print summary results
print("SUMMARY RESULTS")
print("\nOverall Accuracy on Rule Tests:")
for model in ['LR', 'RF', 'NN', 'KAN']:
    acc = results[model]['correct'] / results[model]['total']
    print(f"  {model:5s}: {results[model]['correct']}/{results[model]['total']}␣
    ↳ {acc:.1%}")

print("\nAccuracy by Rule Type:")
print("  " + "-" * 60)
print(f"  {'Rule Category':15s} | {'LR':6s} | {'RF':6s} | {'NN':6s} | {'KAN':
    ↳6s}")
print("  " + "-" * 60)

rule_categories = [
    ('Age (Male)', 'male'),
    ('Age (Female)', 'female'),
    ('Contributions', 'contrib'),
    ('Spouse', 'spouse'),
    ('Absence', 'absent'),
    ('Capital', 'capital'),
    ('In-patient', 'inpatient'),
    ('Out-patient', 'outpatient')
]

for display_name, category in rule_categories:
    row = f"  {display_name:15s} |"
    for model in ['LR', 'RF', 'NN', 'KAN']:
        if category in results[model]['by_rule']:
            acc = results[model]['by_rule'][category]['correct'] /␣
            ↳results[model]['by_rule'][category]['total']
            row += f"  {acc:5.1%} |"
        else:
            row += "  N/A |"
    print(row)

# Visualize results
plot_rule_test_results(results, detailed_results)

# Print failure analysis
print_detailed_failure_analysis(detailed_results)

```

COMPREHENSIVE RULE TESTING

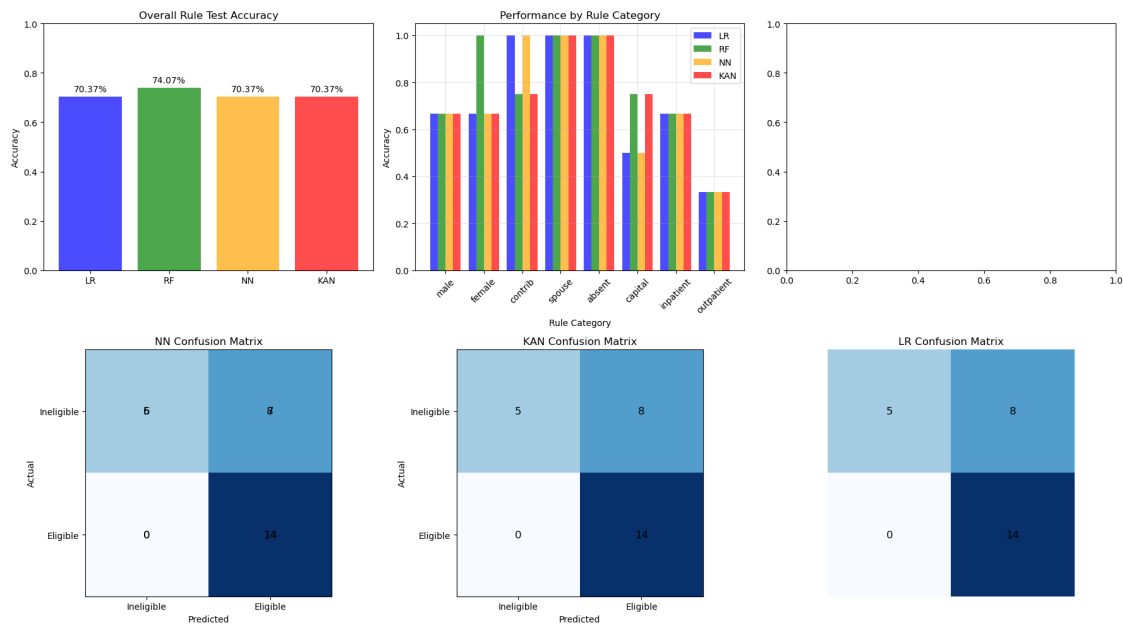
Created 27 test cases covering all rules
SUMMARY RESULTS

Overall Accuracy on Rule Tests:

LR : 19/27 = 70.4%
RF : 20/27 = 74.1%
NN : 19/27 = 70.4%
KAN : 19/27 = 70.4%

Accuracy by Rule Type:

Rule Category	LR	RF	NN	KAN
Age (Male)	66.7%	66.7%	66.7%	66.7%
Age (Female)	66.7%	100.0%	66.7%	66.7%
Contributions	100.0%	75.0%	100.0%	75.0%
Spouse	100.0%	100.0%	100.0%	100.0%
Absence	100.0%	100.0%	100.0%	100.0%
Capital	50.0%	75.0%	50.0%	75.0%
In-patient	66.7%	66.7%	66.7%	66.7%
Out-patient	33.3%	33.3%	33.3%	33.3%



DETAILED FAILURE ANALYSIS

LR Failures (8 total):

male_64_should_fail (prob=0.997)	: Expected INELIGIBLE, Got ELIGIBLE
female_59_should_fail (prob=0.997)	: Expected INELIGIBLE, Got ELIGIBLE
capital_3001_fail (prob=0.994)	: Expected INELIGIBLE, Got ELIGIBLE
capital_5000_fail (prob=0.947)	: Expected INELIGIBLE, Got ELIGIBLE
inpatient_51km_fail (prob=0.998)	: Expected INELIGIBLE, Got ELIGIBLE
outpatient_50km_fail (prob=0.998)	: Expected INELIGIBLE, Got ELIGIBLE
outpatient_30km_fail (prob=0.998)	: Expected INELIGIBLE, Got ELIGIBLE
young_male_few_contrib (prob=0.899)	: Expected INELIGIBLE, Got ELIGIBLE

RF Failures (7 total):

male_64_should_fail (prob=0.930)	: Expected INELIGIBLE, Got ELIGIBLE
contrib_3of5_fail (prob=0.690)	: Expected INELIGIBLE, Got ELIGIBLE
capital_3001_fail (prob=0.729)	: Expected INELIGIBLE, Got ELIGIBLE
inpatient_51km_fail (prob=0.990)	: Expected INELIGIBLE, Got ELIGIBLE
outpatient_50km_fail (prob=0.990)	: Expected INELIGIBLE, Got ELIGIBLE
outpatient_30km_fail (prob=0.990)	: Expected INELIGIBLE, Got ELIGIBLE
young_male_few_contrib (prob=0.836)	: Expected INELIGIBLE, Got ELIGIBLE

NN Failures (8 total):

male_64_should_fail (prob=1.000)	: Expected INELIGIBLE, Got ELIGIBLE
female_59_should_fail (prob=1.000)	: Expected INELIGIBLE, Got ELIGIBLE
capital_3001_fail (prob=0.999)	: Expected INELIGIBLE, Got ELIGIBLE
capital_5000_fail (prob=0.828)	: Expected INELIGIBLE, Got ELIGIBLE
inpatient_51km_fail (prob=1.000)	: Expected INELIGIBLE, Got ELIGIBLE
outpatient_50km_fail (prob=1.000)	: Expected INELIGIBLE, Got ELIGIBLE

```

    outpatient_30km_fail      : Expected INELIGIBLE, Got ELIGIBLE
(prob=1.000)
    young_male_few_contrib    : Expected INELIGIBLE, Got ELIGIBLE
(prob=0.948)

```

KAN Failures (8 total):

```

-----
    male_64_should_fail      : Expected INELIGIBLE, Got ELIGIBLE
(prob=1.000)
    female_59_should_fail    : Expected INELIGIBLE, Got ELIGIBLE
(prob=0.999)
    contrib_3of5_fail        : Expected INELIGIBLE, Got ELIGIBLE
(prob=0.818)
    capital_3001_fail        : Expected INELIGIBLE, Got ELIGIBLE
(prob=0.999)
    inpatient_51km_fail      : Expected INELIGIBLE, Got ELIGIBLE
(prob=1.000)
    outpatient_50km_fail      : Expected INELIGIBLE, Got ELIGIBLE
(prob=1.000)
    outpatient_30km_fail      : Expected INELIGIBLE, Got ELIGIBLE
(prob=1.000)
    young_male_few_contrib    : Expected INELIGIBLE, Got ELIGIBLE
(prob=0.984)

```

10 10. Visualization

```

[21]: fig, axes = plt.subplots(2, 3, figsize=(18, 10))

# Accuracy comparison
models = ['LR', 'RF', 'NN', 'KAN']
accuracies = [lr_acc, rf_acc, nn_acc, kan_acc]
colors = ['blue', 'green', 'orange', 'red']

axes[0, 0].bar(models, accuracies, color=colors, alpha=0.7)
axes[0, 0].set_title('Model Accuracy Comparison')
axes[0, 0].set_ylabel('Accuracy')
for i, acc in enumerate(accuracies):
    axes[0, 0].text(i, acc + 0.005, f'{acc:.3f}', ha='center')

# Rule compliance analysis
lr_coeff = lr.coef_[0]
lr_rules = {
    'age_positive': lr_coeff[feature_cols.index('age')] > 0,
    'spouse_positive': lr_coeff[feature_cols.index('is_spouse')] > 0,
    'absent_negative': lr_coeff[feature_cols.index('is_absent')] < 0,
    'capital_negative': lr_coeff[feature_cols.index('capital_resources')] < 0
}

```

```

lr_compliance = np.mean(list(lr_rules.values()))

rf_top_10 = [f[0] for f in rf_feature_ranking[:10]]
rf_real_count = sum(1 for f in rf_top_10 if f in real_features)
rf_compliance = rf_real_count / 10.0

kan_top_10 = [f[0] for f in kan_feature_ranking[:10]]
kan_real_count = sum(1 for f in kan_top_10 if f in real_features)
kan_compliance = kan_real_count / 10.0

rule_scores = [lr_compliance, rf_compliance, kan_compliance]
axes[0, 1].bar(['LR', 'RF', 'KAN'], rule_scores, color=colors[:3], alpha=0.7)
axes[0, 1].set_title('Rule Compliance Scores')
axes[0, 1].set_ylabel('Compliance Score')
axes[0, 1].set_ylim(0, 1)
for i, score in enumerate(rule_scores):
    axes[0, 1].text(i, score + 0.02, f'{score:.3f}', ha='center')

# Real vs noise feature analysis
real_indices = [i for i, f in enumerate(feature_cols) if f in real_features]
noise_indices = [i for i, f in enumerate(feature_cols) if f in noise_features]

# LR feature focus
lr_importance = np.abs(lr.coef_[0])
lr_real_sum = np.sum(lr_importance[real_indices])
lr_total_sum = np.sum(lr_importance)
lr_focus = lr_real_sum / lr_total_sum

# RF feature focus
rf_real_sum = np.sum(rf.feature_importances_[real_indices])
rf_total_sum = np.sum(rf.feature_importances_)
rf_focus = rf_real_sum / rf_total_sum

# KAN feature focus
kan_real_sum = np.sum(kan_importance[real_indices])
kan_total_sum = np.sum(kan_importance)
kan_focus = kan_real_sum / kan_total_sum

focus_scores = [lr_focus, rf_focus, kan_focus]
axes[0, 2].bar(['LR', 'RF', 'KAN'], focus_scores, color=colors[:3], alpha=0.7)
axes[0, 2].set_title('Real Feature Focus')
axes[0, 2].set_ylabel('Proportion of Importance to Real Features')
axes[0, 2].set_ylim(0, 1)
for i, score in enumerate(focus_scores):
    axes[0, 2].text(i, score + 0.02, f'{score:.2f}', ha='center')

# LR feature importance with color coding

```

```

lr_sorted = sorted(zip(feature_cols, lr_importance), key=lambda x: x[1],
                    ↪reverse=True)
top_10_lr = lr_sorted[:10]
features_lr = [f[0][:12] for f in top_10_lr]
importance_lr = [f[1] for f in top_10_lr]
colors_lr = ['green' if f[0] in real_features else 'red' for f in top_10_lr]

axes[1, 0].barh(range(len(features_lr)), importance_lr, color=colors_lr,
               ↪alpha=0.7)
axes[1, 0].set_yticks(range(len(features_lr)))
axes[1, 0].set_yticklabels(features_lr, fontsize=8)
axes[1, 0].set_title('LR Top 10 Features')
axes[1, 0].set_xlabel('Coefficient Magnitude')

# RF feature importance with color coding
rf_sorted = sorted(rf_feature_ranking, key=lambda x: x[1], reverse=True)
top_10_rf = rf_sorted[:10]
features_rf = [f[0][:12] for f in top_10_rf]
importance_rf = [f[1] for f in top_10_rf]
colors_rf = ['green' if f[0] in real_features else 'red' for f in top_10_rf]

axes[1, 1].barh(range(len(features_rf)), importance_rf, color=colors_rf,
               ↪alpha=0.7)
axes[1, 1].set_yticks(range(len(features_rf)))
axes[1, 1].set_yticklabels(features_rf, fontsize=8)
axes[1, 1].set_title('RF Top 10 Features')
axes[1, 1].set_xlabel('Feature Importance')

# KAN feature importance with color coding
kan_sorted = sorted(kan_feature_ranking, key=lambda x: x[1], reverse=True)
top_10_kan = kan_sorted[:10]
features_kan = [f[0][:12] for f in top_10_kan]
importance_kan = [f[1] for f in top_10_kan]
colors_kan = ['green' if f[0] in real_features else 'red' for f in top_10_kan]

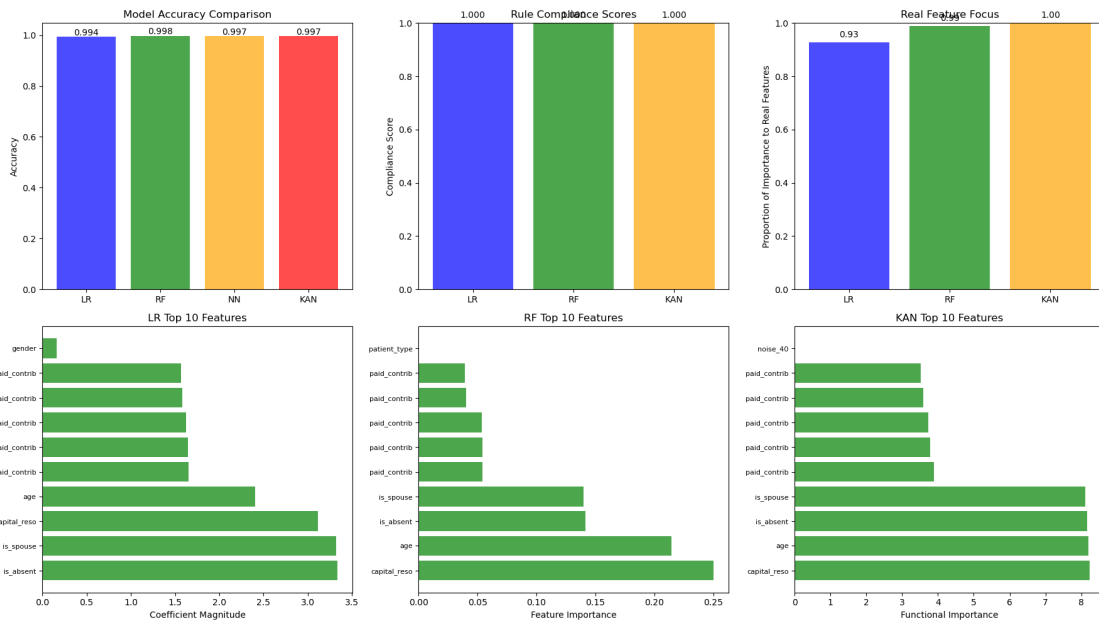
axes[1, 2].barh(range(len(features_kan)), importance_kan, color=colors_kan,
               ↪alpha=0.7)
axes[1, 2].set_yticks(range(len(features_kan)))
axes[1, 2].set_yticklabels(features_kan, fontsize=8)
axes[1, 2].set_title('KAN Top 10 Features')
axes[1, 2].set_xlabel('Functional Importance')

plt.tight_layout()
plt.show()

# Summary statistics
print(f"LR Real Feature Focus: {lr_focus:.3f}")

```

```
print(f"RF Real Feature Focus: {rf_focus:.3f}")
print(f"KAN Real Feature Focus: {kan_focus:.3f}")
```



LR Real Feature Focus: 0.927
RF Real Feature Focus: 0.988
KAN Real Feature Focus: 1.000